

Objective: Our goal is to find how quantity of each item in each province, changes from each month and year. During milestone 1, we try to finish all the preprocessing steps and draw necessary graphs for insights. We will try to build the model and make predictions in milestone 2 and later stages.

#Import Files

```
from google.colab import files
uploaded = files.upload()
```

No file chosen

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving item.csv to item.csv

Saving train.csv to train.csv

#Read data

```
import pandas as pd
import numpy as np
df_item = pd.read_csv('item.csv')
df_train = pd.read_csv('train.csv')
print(df_item)
print(df_train)
```

```

[ ] item_id brand_id brandfamily_id package_id info_id
0      1  [6, 14]           1           1    BHS
1      6  [6, 19]           2           1    BHM
2     54  [6, 14]           1           1    BHS
3     15  [6, 19]           2           1    BHM
4    108  [19]             3           1    MQB
..     ...      ...           ...         ...
113    78  [6, 8]           4           4    BZMN
114     5  [11]            5           1    HGO
115    83  [19]            3           1    MQB
116    87  [19]            3           2    MQB
117    12  [19]            3           4    MQB

```

[118 rows x 5 columns]

```

      datetime  order_id  user_id  item_id  province_id  city_id  quantity
0    2017-01-02     8806     404         1           2       19    8.298021
1    2017-01-02    22552     404         1           2       19    8.298021
2    2017-01-02     494     489         6           1        0    0.376088
3    2017-01-02    62721      28        54           2       19    1.488643
4    2017-01-02    24790    231        15           0       10    0.748268
...         ...      ...      ...      ...         ...      ...      ...
145190 2020-12-29    72305    390         18           0       10    0.733838
145191 2020-12-30    55224    321         27           1        3    0.792584
145192 2020-12-31    75987    326        105           0        4    0.009584
145193 2020-12-31    27911    326        105           0        4    0.015935
145194 2020-12-31    55625    326         75           0        4    0.069767

```

[145195 rows x 7 columns]

We find that our two tables have a common column called id, we need to combine them in order to continue the analysis.

```
#Data integrity
df_wine = pd.merge(df_item,df_train, on = 'item_id')
df_wine
```

	item_id	brand_id	brandfamily_id	package_id	info_id	datetime	order_id
<b>0</b>	1	[6, 14]	1	1	BHS	2017-01-02	8806
<b>1</b>	1	[6, 14]	1	1	BHS	2017-01-02	22552
<b>2</b>	1	[6, 14]	1	1	BHS	2017-01-06	54066
<b>3</b>	1	[6, 14]	1	1	BHS	2017-01-07	67828
<b>4</b>	1	[6, 14]	1	1	BHS	2017-01-10	4205
...	...	...	...	...	...	...	...
<b>145190</b>	12	[19]	3	4	MQB	2020-09-21	52082
<b>145191</b>	12	[19]	3	4	MQB	2020-09-26	11453
						2020-09-26	

The join result get a table with 145195 rows and 11 columns, which is good

```
#Explode composite values
from ast import literal_eval
df_wine['brand_id'] = df_wine['brand_id'].astype(str)
df_wine['brand_id'] = df_wine['brand_id'].apply(literal_eval)
df_wine = df_wine.explode('brand_id')
df_wine
```

item_id	brand_id	brandfamily_id	package_id	info_id	datetime	order_id
0	1	6	1	1	BHS	2017-01-02
					2017-01-	8806

We find that brand\_id has composite items(it is a list), wh need to change its types to numeric.  
Explode the lists to get more rows(217723)

#Data Cleaning

#Check Data Type

df\_wine.dtypes

```

item_id          int64
brand_id         object
brandfamily_id   int64
package_id       int64
info_id         object
datetime        object
order_id        int64
user_id         int64
province_id      int64
city_id         int64
quantity        float64
dtype: object

```

We check the types find that some types are unclear, they are object which are hard to handle

from datetime import datetime

#Convert types

df\_wine['brand\_id'] = df\_wine['brand\_id'].astype(int)

df\_wine['info\_id'] = df\_wine['info\_id'].astype('category')

df\_wine['datetime'] = pd.to\_datetime(df\_wine['datetime'], format='%Y-%m-%d')

df\_wine.dtypes

```

item_id          int64
brand_id         int64
brandfamily_id   int64
package_id       int64
info_id         category
datetime        datetime64[ns]
order_id        int64
user_id         int64
province_id      int64
city_id         int64
quantity        float64
dtype: object

```

# Convert string id to numeric

df\_wine['info\_id'] = df\_wine['info\_id'].cat.codes

df\_wine['info\_id'] = df\_wine['info\_id'].astype(int)

df\_wine.dtypes

```
df_wine.dtypes
```

```

item_id          int64
brand_id         int64
brandfamily_id   int64
package_id       int64
info_id          int64
datetime         datetime64[ns]
order_id         int64
user_id          int64
province_id      int64
city_id          int64
quantity         float64
dtype: object

```

We try to make sure every id is int. Convert brand\_id to int. Cast info\_id into category then convert into int. Convert the datetime from object to datetime object.

```
#Cleaning NA values
```

```
df_wine.dropna(axis = 0, how = "any")
df_wine
```

	item_id	brand_id	brandfamily_id	package_id	info_id	datetime	order_id
0	1	6	1	1	3	2017-01-02	8806
0	1	14	1	1	3	2017-01-02	8806
1	1	6	1	1	3	2017-01-02	22552
1	1	14	1	1	3	2017-01-02	22552
2	1	6	1	1	3	2017-01-06	54066
...	...	...	...	...	...	...	...
145190	12	19	3	4	11	2020-09-21	52082
145191	12	19	3	4	11	2020-09-26	11453
...	...	...	...	...	...	...	...

We check NA through all rows and columns. We find that there are no missing values at all, which means the table is good itself.

```
#Check negative values
```

```
print("The unusual values for item id:",df_wine[ df_wine.iloc[:, 0] < 0 ])
```

```

print("The unusual values for brand id:",df_wine[ df_wine.iloc[: , 1] < 0 ])
print("The unusual values for brandfamily id:",df_wine[ df_wine.iloc[: , 2] < 0 ])
print("The unusual values for package id:",df_wine[ df_wine.iloc[: , 3] < 0 ])
print("The unusual values for order id:",df_wine[ df_wine.iloc[: , 6] < 0 ])
print("The unusual values for user id:",df_wine[ df_wine.iloc[: , 7] < 0 ])
print("The unusual values for province id:",df_wine[ df_wine.iloc[: , 8] < 0 ])
print("The unusual values for city id:",df_wine[ df_wine.iloc[: , 9] < 0 ])
print("The unusual values for quantity :",df_wine[ df_wine.iloc[: , 10] < 0 ])

```

The unusual values for item id: Empty DataFrame

Columns: [item\_id, brand\_id, brandfamily\_id, package\_id, info\_id, datetime, order\_id,  
Index: []

The unusual values for brand id: Empty DataFrame

Columns: [item\_id, brand\_id, brandfamily\_id, package\_id, info\_id, datetime, order\_id,  
Index: []

The unusual values for brandfamily id: Empty DataFrame

Columns: [item\_id, brand\_id, brandfamily\_id, package\_id, info\_id, datetime, order\_id,  
Index: []

The unusual values for package id: Empty DataFrame

Columns: [item\_id, brand\_id, brandfamily\_id, package\_id, info\_id, datetime, order\_id,  
Index: []

The unusual values for order id: Empty DataFrame

Columns: [item\_id, brand\_id, brandfamily\_id, package\_id, info\_id, datetime, order\_id,  
Index: []

The unusual values for user id: Empty DataFrame

Columns: [item\_id, brand\_id, brandfamily\_id, package\_id, info\_id, datetime, order\_id,  
Index: []

The unusual values for province id: Empty DataFrame

Columns: [item\_id, brand\_id, brandfamily\_id, package\_id, info\_id, datetime, order\_id,  
Index: []

The unusual values for city id: Empty DataFrame

Columns: [item\_id, brand\_id, brandfamily\_id, package\_id, info\_id, datetime, order\_id,  
Index: []

The unusual values for quantity : Empty DataFrame

Columns: [item\_id, brand\_id, brandfamily\_id, package\_id, info\_id, datetime, order\_id,  
Index: []



We check whether there are negative values in every id columns and also quantity columns(not alllowed, should be removed as they are unusual values). The good news is that there is no negative values through those columns.

#split date

```
df_wine.insert(0, 'month', df_wine['datetime'].dt.month)
```

```
df_wine.insert(0, 'year', df_wine['datetime'].dt.year)
```

```
df_wine.insert(0, 'date', pd.to_datetime(df_wine[['year', 'month']].assign(DAY=1)))
```

```
df_wine = df_wine.drop('datetime', axis=1)
```

```
df_wine = df_wine.drop('year', axis=1)
```

```
df_wine = df_wine.drop('month', axis=1)
```

```
df_wine
```

	date	item_id	brand_id	brandfamily_id	package_id	info_id	order_id	u:
<b>0</b>	2017-01-01	1	6	1	1	3	8806	
<b>0</b>	2017-01-01	1	14	1	1	3	8806	
<b>1</b>	2017-01-01	1	6	1	1	3	22552	
<b>1</b>	2017-01-01	1	14	1	1	3	22552	
<b>2</b>	2017-01-01	1	6	1	1	3	54066	
...	...	...	...	...	...	...	...	
<b>145190</b>	2020-09-01	12	19	3	4	11	52082	
~~~~~								

We split the date into year and month, and combine again, as our target is to predict through year and month(day is not important)

```

from sklearn.preprocessing import StandardScaler
from sklearn import preprocessing
print(df_wine.iloc[:,1:10])
data_scaled = pd.DataFrame(preprocessing.scale(df_wine.iloc[:,1:10]),columns = df_wine.ilo
df_array = np.array(df_wine.iloc[:,1:10])
X_std = StandardScaler().fit_transform(df_array)
from sklearn.decomposition import PCA
pca = PCA(n_components = 4)
principalComponents = pca.fit_transform(X_std)
print("PCA Variance",pca.explained_variance_)
print("PCA Variance",pca.explained_variance_ratio_)
print("Cumulative PCA Variance",np.cumsum(pca.explained_variance_ratio_))
principalDf = pd.DataFrame(data = principalComponents)
principalDf
result = pd.DataFrame(pca.components_,columns=data_scaled.columns)
print(result.abs().max(axis=1))
result.abs()

```

	item_id	brand_id	brandfamily_id	...	user_id	province_id	city_id
0	1	6	1	...	404	2	19
0	1	14	1	...	404	2	19
1	1	6	1	...	404	2	19
1	1	14	1	...	404	2	19
2	1	6	1	...	395	0	10
...	...	...	...	...	...	...	...
145190	12	19	3	...	18	0	10
145191	12	19	3	...	450	0	10
145192	12	19	3	...	262	0	10
145193	12	19	3	...	248	0	10
145194	12	19	3	...	24	0	10

[217723 rows x 9 columns]

PCA Variance [2.195108 1.41419053 1.14127356 1.01680283]

PCA Variance [0.24389977 0.15713156 0.12680759 0.11297757]

Cumulative PCA Variance [0.24389977 0.40103133 0.52783892 0.64081649]

0 0.491384

1 0.537006

2 0.565828

3 0.775233

dtype: float64

item id brand id brandfamily id package id info id order id user id pr

We know province\_id, brandfamily\_id, brand\_id, item\_id are most important features from PCA .

As our main target is to predict each item in each province through each month and year, we will try to remove unnecessary columns.

#dimension reduction

```
df_wine = df_wine.drop(['package_id', 'info_id', 'order_id', 'user_id', 'city_id'], axis=1)
df_wine
```

	date	item_id	brand_id	brandfamily_id	province_id	quantity
0	2017-01-01	1	6	1	2	8.298021
0	2017-01-01	1	14	1	2	8.298021
1	2017-01-01	1	6	1	2	8.298021
1	2017-01-01	1	14	1	2	8.298021
2	2017-01-01	1	6	1	0	0.377183
...	...	...	...	...	...	...
145190	2020-09-01	12	19	3	0	3.643725
145191	2020-09-01	12	19	3	0	1.457490
145192	2020-09-01	12	19	3	0	2.914980
145193	2020-09-01	12	19	3	0	1.457490
145194	2020-09-01	12	19	3	0	0.728745

217723 rows x 6 columns

We remove the columns 'package\_id', 'info\_id', 'order\_id', 'user\_id', 'city\_id' from the tables. They are insensitive to analysis and also not our anylysis focus in the prediction.

```
df_wine = df_wine.groupby(['date', 'province_id', 'item_id', 'brand_id', 'brandfamily_id'])[
df_wine
```

	date	province_id	item_id	brand_id	brandfamily_id	quantity
<b>0</b>	2017-01-01	0	1	6	1	421.313155
<b>1</b>	2017-01-01	0	1	14	1	421.313155
<b>2</b>	2017-01-01	0	2	11	5	5.654450
<b>3</b>	2017-01-01	0	15	6	2	117.478060
<b>4</b>	2017-01-01	0	15	19	2	117.478060
...	...	...	...	...	...	...
<b>8579</b>	2020-12-01	2	109	6	2	46.565716
<b>8580</b>	2020-12-01	2	109	19	2	46.565716
<b>8581</b>	2020-12-01	2	110	13	9	128.102211
<b>8582</b>	2020-12-01	2	111	6	2	23.734884
<b>8583</b>	2020-12-01	2	111	19	2	23.734884

8584 rows × 6 columns

We aggregate the quantity through all remaning columns to get how each item in each province varies through each month and year.

```
import seaborn as sns
```

```
#Visulize Outliers
```

```
sns.scatterplot(x='date', y='quantity', data=df_wine)
```

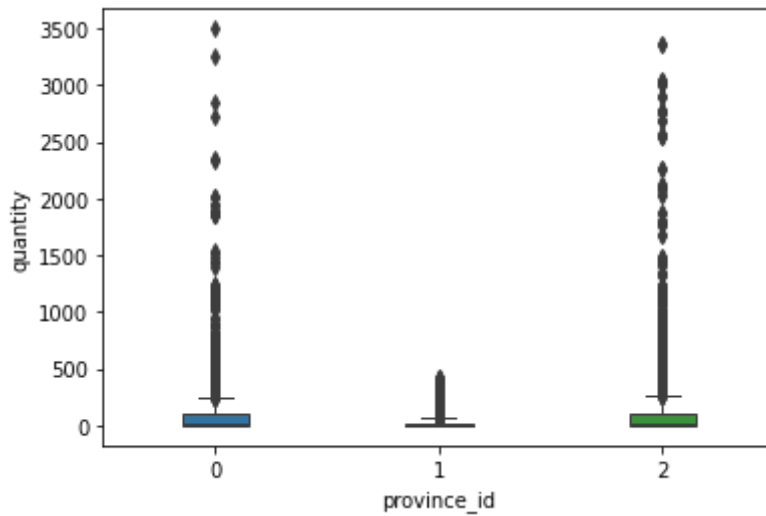


```
<matplotlib.axes._subplots.AxesSubplot at 0x7fdd4479ea50>
```



```
sns.boxplot(x = 'province_id', y='quantity', data=df_wine , width = 0.3, linewidth = 1)
```


```
<matplotlib.axes._subplots.AxesSubplot at 0x7fdd447212d0>
```



```
# Violin plots
```

```
sns.catplot(x = 'province_id', y='quantity', data=df_wine, kind="violin")
```

```
<seaborn.axisgrid.FacetGrid at 0x7fdd446ac050>
```



Obviously the graphs look not normal and has a lot of outliers and one side, we need to handle them immediately

```
Q1 = df_wine['quantity'].quantile(0.25)
```

```
Q3 = df_wine['quantity'].quantile(0.75)
```

```
IQR = Q3 - Q1
```

```
index = df_wine[(df_wine['quantity'] < (Q1 - 1.5 * IQR)) | (df_wine['quantity'] > (Q3 + 1.5 * IQR))]
```

```
index = df_wine[(df_wine['quantity'] < (Q1 - 1.5 * IQR)) | (df_wine['quantity'] > (Q3 + 1.5 * IQR)]
df_wine.drop(index, inplace=True)
df_wine
```


	date	province_id	item_id	brand_id	brandfamily_id	quantity
2	2017-01-01	0	2	11	5	5.654450
3	2017-01-01	0	15	6	2	117.478060
4	2017-01-01	0	15	19	2	117.478060
5	2017-01-01	0	18	19	3	57.459522
6	2017-01-01	0	19	11	5	43.250291
...	...	...	...	...	...	...
8579	2020-12-01	2	109	6	2	46.565716
8580	2020-12-01	2	109	19	2	46.565716
8581	2020-12-01	2	110	13	9	128.102211
8582	2020-12-01	2	111	6	2	23.734884
8583	2020-12-01	2	111	19	2	23.734884

7451 rows × 6 columns

Through outliers drop we remove around 1100 rows through the IQR rule. We then check through visualization again.


```
sns.scatterplot(x='date', y='quantity', data=df_wine)
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7fdd44619ed0>




```
sns.boxplot(x = 'province_id', y='quantity', data=df_wine , width = 0.3, linewidth = 1)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fdd44629690>
```



```
sns.catplot(x = 'province_id', y='quantity', data=df_wine, kind="violin")
```

```
<seaborn.axisgrid.FacetGrid at 0x7fdd446a7a10>
```




We still get some quantities which are large. However, the overall performance is much better. There are much fewer values beyond.

```
#Data Visualization
```

```
sns.displot(df_wine, x='quantity', kind = 'kde',bw_adjust = 5)
```

```
<seaborn.axisgrid.FacetGrid at 0x7fdd4445eb50>
```




From density plot of quantity, we see that most quantity are small and near 0, only a small part are large

```
|      /      \
```

```
sns.histplot(df_wine, x='item_id', binwidth = 3, bins=df_wine['item_id'].nunique())
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fdd4441e490>
```



We plot the histograms through id and find there is huge difference between each item.

```
sns.catplot(x='province_id', y='quantity', kind="bar", data=df_wine)
```

```
<seaborn.axisgrid.FacetGrid at 0x7fdd443873d0>
```




Use bar chart we find that province 0 has most quantities, followed by province 2. Province 1 has least quantity.



```
sns.heatmap(df_wine.corr(), annot=True, cmap="Blues")
```


```
<matplotlib.axes._subplots.AxesSubplot at 0x7fdd440cfe10>
```



From the heatmap, we find that there is not much direct relationship between each predictor (and also quantity). We can say they are almost independent with each other.

```
sns.lineplot(data = df_wine, x='date', y='quantity')
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fdd43f2d050>
```



We have a sense of how overall quantity varies with month and year. It drops sharply at

• `df_wine.to_csv('My Drive/Colab Notebooks/IE7275/Project/wine.csv')`

```
from google.colab import drive
```

```
drive.mount('drive')
```

```
#df_wine.to_csv('My Drive/Colab Notebooks/IE7275/Project/wine.csv')
```

```
df_wine.to_csv('/content/drive/MyDrive/Colab Notebooks/IE7275/Project/wine.csv')
```

Drive already mounted at drive; to attempt to forcibly remount, call `drive.mount("drive")`



Save the new table to google drive and download in convenience of future use.